

Classification of Sentinel-2 Multispectral Images using Deep Neural Network in R-H2O (Windows 10)

Zia Ahmed, PhD, Research Associate Professor, RENEW, University at Buffalo

March 22, 2018

This tutorial will show how to implement [Deep Neural Network](#) for pixel based [supervised classification](#) of [Sentinel-2 multispectral images](#) using [H2O](#) package in [R](#).

[H2O](#) is an open source, in-memory, distributed, fast, and scalable machine learning and predictive analytics platform that allows you to build machine learning models on big data and provides easy productionalization of those models in an enterprise environment. It's core code is written in Java and can read data in parallel from a distributed cluster and also from local cluster. H2O allows access to all the capabilities of H2O from an external program or script via JSON over HTTP. The Rest API is used by H2O's [web interface \(Flow UI\)](#), [R binding \(H2O-R\)](#), and [Python binding \(H2O-Python\)](#). Requirement and installation steps in R can be found here [here](#).

We will use the Deep Neural Network algorithm using [H2O](#) package in [R](#) for image classification. First, we will split "point_data" into a training set (75% of the data), a validation set (12%) and a test set (13%) data. The validation data set will be used to optimize the model parameters during training process. The model's performance will be tested with the data set and then we will predict landuse classes on grid data set. The point and grid data can be download as [rar](#), [7z](#) and [zip](#) format.

Tuning and Optimizations parameters:

- Four hidden layers with 200 neurons and Rectifier Linear (ReLU) as a activation function of neurons.
- The default stochastic gradient descent function will be used to optimize different objective functions and to minimize training loss.
- To reduce the generalization error and the risk of over-fitting of the model, we will use set low values for L1 and L2 regularizations.
- The model will be cross validated with 10 folds with stratified sampling
- The model will be run with 100 epochs.

More details of Tuning and Optimizations parameters of H2O Deep Neural Network for supervised classification can be found [here](#)

```
start_time <- Sys.time()
```

Load packages

```
library(rgdal) # spatial data processing
library(raster) # raster processing
library(plyr) # data manipulation
```

```
library(dplyr) # data manipulation
library(RStoolbox) # plotting spatial data
library(ggplot2) # plotting
library(RColorBrewer)
library(sp)
```

Set working directory

```
setwd("F:\\My_GitHub\\DNN_H2O_R")
```

Load point and grid data

```
point<-read.csv("point_data.csv", header = T)
grid<-read.csv("grid_data.csv", header = T)
```

Creat data frames

```
point.data<-cbind(point[c(4:13)],Class=point$Class)
grid.data<-grid[c(4:13)]
grid.xy<-grid[c(3,1:2)]
```

Install H2O

```
#install.packages("h2o")
```

Start and Initialize H2O local cluster

```
library(h2o)
localH2o <- h2o.init(nthreads = -1, max_mem_size = "50G")
```

Import data to H2O cluster

```
df<- as.h2o(point.data)
grid<- as.h2o(grid.data)
```

Split data into train, validation and test dataset

```
splits <- h2o.splitFrame(df, c(0.75,0.125), seed=1234)
train  <- h2o.assign(splits[[1]], "train.hex") # 75%
valid  <- h2o.assign(splits[[2]], "valid.hex") # 12%
test   <- h2o.assign(splits[[3]], "test.hex")  # 13%
```

Create response and features data sets

```
y <- "Class"
x <- setdiff(names(train), y)
```

Deep Learning Model

```

d1_model <- h2o.deeplearning(
  model_id="Deep_Learning",
  training_frame=train,
  validation_frame=valid,
  x=x,
  y=y,
  standardize=TRUE,
  score_training_samples=0,
  activation = "RectifierWithDropout",
  score_each_iteration = TRUE,
  hidden = c(200,200,200,200),
  hidden_dropout_ratios=c(0.2,0.1,0.1,0),
  stopping_tolerance = 0.001,
  epochs=100,
  adaptive_rate=TRUE,
  l1=1e-6,
  l2=1e-6,
  max_w2=10,
  # Destination id for this model
  # Id of the training data frame
  # Id of the validation data frame
  # a vector predictor variable
  # name of reponse variables
  # standardize the data
  # training set samples for scoring (0 for all)
  # Activation function
  # 4 hidden layers, each of 200 neurons
  # for improve generalization
  # tolerance for metric-based stopping criterion
  # the dataset should be iterated (streamed)
  # manually tuned Learning rate
  # L1/L2 regularization, improve generalization
  # helps stability for Rectifier

```

```

nolds=10,                                # Number of folds for K-fold cross-validation
fold_assignment="Stratified",            # Cross-validation fold assignment scheme
keep_cross_validation_fold_assignment = TRUE,
seed=125,
reproducible = TRUE,
variable_importances=T
)

```

Model Summary

```

#summary(dl_model)
#capture.output(print(summary(dl_model)),file = "DL_summary_model_01.txt")

```

Mean error

```

h2o.mean_per_class_error(dl_model, train = TRUE, valid = TRUE, xval = TRUE)

```

```

##      train      valid      xval
## 0.005951827 0.008541536 0.009924872

```

Scoring history

```

scoring_history<-dl_model@model$scoring_history
#write.csv(scoring_history, "scoring_history_model_02.csv")


```

Plot the classification error

```

plot(dl_model,
      timestep = "epochs",
      metric = "classification_error")

```

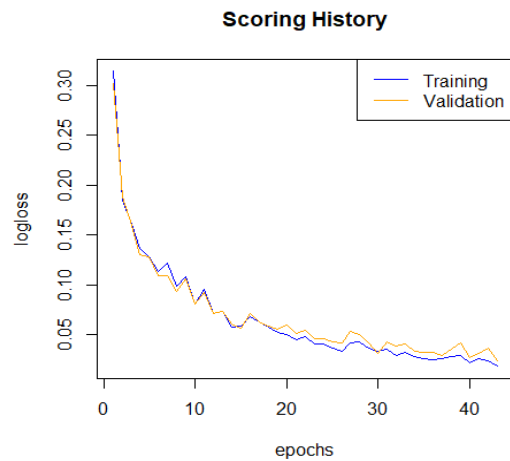


Plot logloss

```


plot(dl_model,
      timestep = "epochs",
      metric = "logloss")

```




Plot RMSE

```
plot(dl_model,
      timestep = "epochs",
      metric = "rmse")
```



Cross-validation Error

```
# Get the CV models from the deeplearning model object` object
cv_models <- sapply(dl_model@model$cross_validation_models,
                    function(i) h2o.getModel(i$name))
# Plot the scoring history over time
plot(cv_models[[1]],
      timestep = "epochs",
      metric = "classification_error")
```



Cross validation result

```
print(dl_model@model$cross_validation_metrics_summary%>%[,c(1,2)])
```

```
##               mean          sd
## accuracy      0.989942  0.00306049
## err           0.010057977 0.00306049
## err_count      18.2      5.681549
## logloss       0.032726314 0.01081141
## max_per_class_error 0.02872035 0.0097002955
## mean_per_class_accuracy 0.99011856 0.0029896488
## mean_per_class_error 0.009881463 0.0029896488
## mse           0.008144272 0.0023309013
## r2            0.99390364 0.0017807437
## rmse          0.088454194 0.012651616
```

```
#capture.output(print(dl_model@model$cross_validation_metrics_summary%>%[,c(1,2)]),file = "DL_CV_model_01.txt")
```

Model performance with Test data set

Compare the training error with the validation and test set errors

```
h2o.performance(dl_model, newdata=train) ## full train data
```

```
## H2OMultinomialMetrics: deeplearning
##
## Test Set Metrics:
## =====
##
## MSE: (Extract with `h2o.rmse`) 0.004795433
## RMSE: (Extract with `h2o.rmse`) 0.06924907
## Logloss: (Extract with `h2o.logloss`) 0.01829286
## Mean Per-Class Error: 0.005951827
## Confusion Matrix: Extract with `h2o.confusionMatrix(<model>, <data>)`
## =====
## Confusion Matrix: Row labels: Actual class; Column labels: Predicted class
##
##      Class_1 Class_2 Class_3 Class_4 Class_5 Error Rate
## Class_1    4150      15      1      0      0 0.0038 = 16 / 4,166
## Class_2      48    3249     23      0      0 0.0214 = 71 / 3,320
## Class_3     11      0    6102      2      0 0.0021 = 13 / 6,115
## Class_4      0      0      9    3730      0 0.0024 = 9 / 3,739
## Class_5      0      0      0      0    694 0.0000 = 0 / 694
```

```
## Totals      4209      3264      6135      3732      694 0.0060 = 109 / 18,034
##
## Hit Ratio Table: Extract with `h2o.hit_ratio_table(<model>, <data>)`
## =====
## Top-5 Hit Ratios:
##   k hit_ratio
## 1 1 0.993956
## 2 2 1.000000
## 3 3 1.000000
## 4 4 1.000000
## 5 5 1.000000
```

h2o.performance(dl_model, newdata=valid) ## full validation data

```
## H2OMultinomialMetrics: deeplearning
##
## Test Set Metrics:
## =====
##
## MSE: (Extract with `h2o.mse`) 0.00687076
## RMSE: (Extract with `h2o.rmse`) 0.08289005
## Logloss: (Extract with `h2o.logloss`) 0.02419423
## Mean Per-Class Error: 0.008541536
## Confusion Matrix: Extract with `h2o.confusionMatrix(<model>, <data>)`
## =====
## Confusion Matrix: Row labels: Actual class; Column labels: Predicted class
##      Class_1 Class_2 Class_3 Class_4 Class_5 Error      Rate
## Class_1      657      3      0      0      0 0.0045 =   3 / 660
## Class_2       8     529      7      0      0 0.0276 =  15 / 544
## Class_3       3      0     999      1      0 0.0040 =   4 / 1,003
## Class_4       0      0      4     602      0 0.0066 =   4 / 606
## Class_5       0      0      0      0     111 0.0000 =   0 / 111
## Totals      668     532     1010     603     111 0.0089 =  26 / 2,924
##
## Hit Ratio Table: Extract with `h2o.hit_ratio_table(<model>, <data>)`
## =====
## Top-5 Hit Ratios:
##   k hit_ratio
## 1 1 0.991108
## 2 2 1.000000
## 3 3 1.000000
## 4 4 1.000000
## 5 5 1.000000
```

h2o.performance(dl_model, newdata=test) ## full test data

```
## H2OMultinomialMetrics: deeplearning
##
## Test Set Metrics:
## =====
##
## MSE: (Extract with `h2o.mse`) 0.00642109
## RMSE: (Extract with `h2o.rmse`) 0.0801317
## Logloss: (Extract with `h2o.logloss`) 0.02786502
## Mean Per-Class Error: 0.007264487
## Confusion Matrix: Extract with `h2o.confusionMatrix(<model>, <data>)`
## =====
## Confusion Matrix: Row labels: Actual class; Column labels: Predicted class
##      Class_1 Class_2 Class_3 Class_4 Class_5 Error      Rate
## Class_1      704      3      1      0      0 0.0056 =   4 / 708
## Class_2       8     553      4      0      0 0.0212 =  12 / 565
```

```
## Class_3      3      0      976      0      0 0.0031 = 3 / 979
## Class_4      0      0      4      624      0 0.0064 = 4 / 628
## Class_5      0      0      0      0      107 0.0000 = 0 / 107
## Totals      715     556     985     624     107 0.0077 = 23 / 2,987
##
## Hit Ratio Table: Extract with `h2o.hit_ratio_table(<model>, <data>)`
## =====
## Top-5 Hit Ratios:
##   k hit_ratio
## 1 1 0.992300
## 2 2 1.000000
## 3 3 1.000000
## 4 4 1.000000
## 5 5 1.000000

#capture.output(print(h2o.performance(dl_model, test)), file =
"test_data_model_01.txt")
```

Confusion matrix

```
train.cf<-h2o.confusionMatrix(dl_model)
print(train.cf)
valid.cf<-h2o.confusionMatrix(dl_model, valid=TRUE)
print(valid.cf)
test.cf<-h2o.confusionMatrix(dl_model, test)
print(test.cf)
#write.csv(train.cf, "CFM_train_model_01.csv")
#write.csv(valid.cf, "CFM_valid_model_01.csv")
#write.csv(test.cf, "CFM_test_model_01.csv")
```

Grid Prediction

```
g.predict = as.data.frame(h2o.predict(object = dl_model, newdata = grid))
```

Stop h2o cluster

```
h2o.shutdown(prompt=FALSE)
```

```
## [1] TRUE
```

Extract Prediction Class

```
grid.xy$Class<-g.predict$predict
str(grid.xy)

## 'data.frame':   35700 obs. of  4 variables:
## $ ID   : int  1 2 3 4 5 6 7 8 9 10 ...
## $ x    : int  677775 677785 677795 677805 677815 677825 677835 677845 677855 677865 ...
## $ y    : int  4764065 4764065 4764065 4764065 4764065 4764065 4764065 4764065 4764065 4764065 ...
## $ Class: Factor w/ 5 levels "Class_1","Class_2",...: 1 1 1 1 1 1 3 1 2 1 ...

grid.xy.na<-na.omit(grid.xy)
```

Join Class Id Column

```
ID<-read.csv("Landuse_ID_h2o.csv", header=TRUE)
new.grid<-join(grid.xy.na, ID, by="Class", type="inner")
#write.csv(new.grid, "Predicted_Landuse_Class.csv")
```

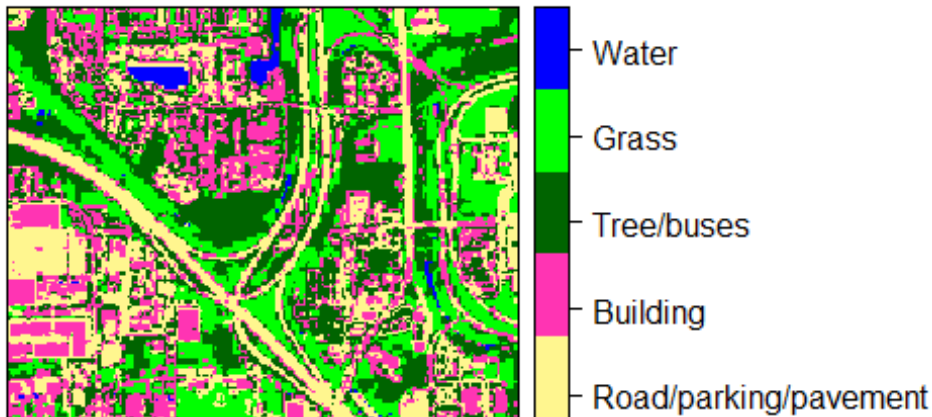
Convert to raster and write

```
x<-SpatialPointsDataFrame(as.data.frame(new.grid)[, c("x", "y")], data = new.grid)
r <- rasterFromXYZ(as.data.frame(x)[, c("x", "y", "Class_ID")])
#writeRaster(r, "predicted_Landuse.tif", "GTiff", overwrite=TRUE)
```

Plot and Save as a tiff file

```
myPalette <- colorRampPalette(c("khaki1", "maroon1", "darkgreen", "green", "blue"))
lu<-spplot(r, "Class_ID", main="Landuse Classes",
  colorkey = list(space="right", tick.number=1, height=1, width=1.5,
    labels = list(at = seq(1,4.8, length=5), cex=1.0,
      lab = c("Road/parking/pavement", "Building", "Tree/buses", "Grass", "Water")),
    col.regions=myPalette, cut=4)
lu
```

Landuse Classes



```
# Save as tif file
windows(width=4, height=4)
tiff( file="FIGURE_Landuse_Class.tif",
  width=4,
  height=4,
  units = "in",
  pointsize = 12,
  res=600,
  restoreConsole = T,
  compression = "lzw",
  bg="transparent")
print(lu)
dev.off()

## png
## 2
```


Run time

```
end_time <- Sys.time()  
end_time - start_time
```

```
## Time difference of 30.42456 mins
```

Conclusions

This simple pixel-based satellite image classification algorithm with deep neural network with H2O-R able to identify urban objects with very high accuracy. It may be use full for landuse classification for urban environment monitoring as well as planning purpose. Also, may use full for agricultural landuse classification.